



# Agent Shield

The token-trust layer for autonomous AI agents

*Agents that decide independently — and refuse  
scams on their own.*

● LIVE

Pharos Atlantic · 0x5ad1...aCe7 · verified

## WHY NOW

# Three curves just crossed.

- 01 AI agents now hold wallets and move funds on their own
- 02 On-chain scams scale faster than any human can vet
- 03 No shared safety layer exists for agents — yet



The moment agents transact autonomously, they need a trust layer. That moment is now.



## THE PROBLEM

# Autonomous agents, autonomous losses.



**A honeypot looks identical to a real token**

until the agent tries to sell — and can't.



**An agent can't pause to feel suspicious**

it acts in milliseconds, with real funds.



**Every agent re-checking every token**

wastes the whole network's effort, forever.

• **WARN** \$3.3M vanished in the Squid Game token alone — sells were silently blocked.

THE INSIGHT

Trust shouldn't be  
recomputed by every  
agent. It should be **written  
once, and read by all.**

Scanners check a token for one user. Agent Shield turns the verdict into a permanent, shared, on-chain asset.

THE SOLUTION

# An immune system for the agent economy

## SCREEN

GoPlus intel + on-chain  
heuristics

→ a 0-100 score

## ATTEST

Verdict on-chain

ALLOW

WARN

BLOCK

Permanent. No admin key.

## REUSE

Every other agent reads it in  
one free call

One check protects thousands.

Unknown ≠ safe.

# Off-chain intelligence, on-chain memory.

The loop fills itself as agents  
protect themselves.





# The agent decides — by itself

- 1 “Should I buy token X?”
- 2 Read registry: any verdict?
- 3 If unknown → screen it live
- 4 Score 5 · **BLOCK** · honeypot
- 5 **Agent refuses the trade**

• **BLOCK**

verdict

5 / 100

HONEYPOT · SELL\_RESTRICTED

*No human in the loop. The skill teaches any agent to do this.*

# Security that pays for itself.



## WHO PAYS

agents / protocols / users needing a token screened

## UNIT ECONOMICS

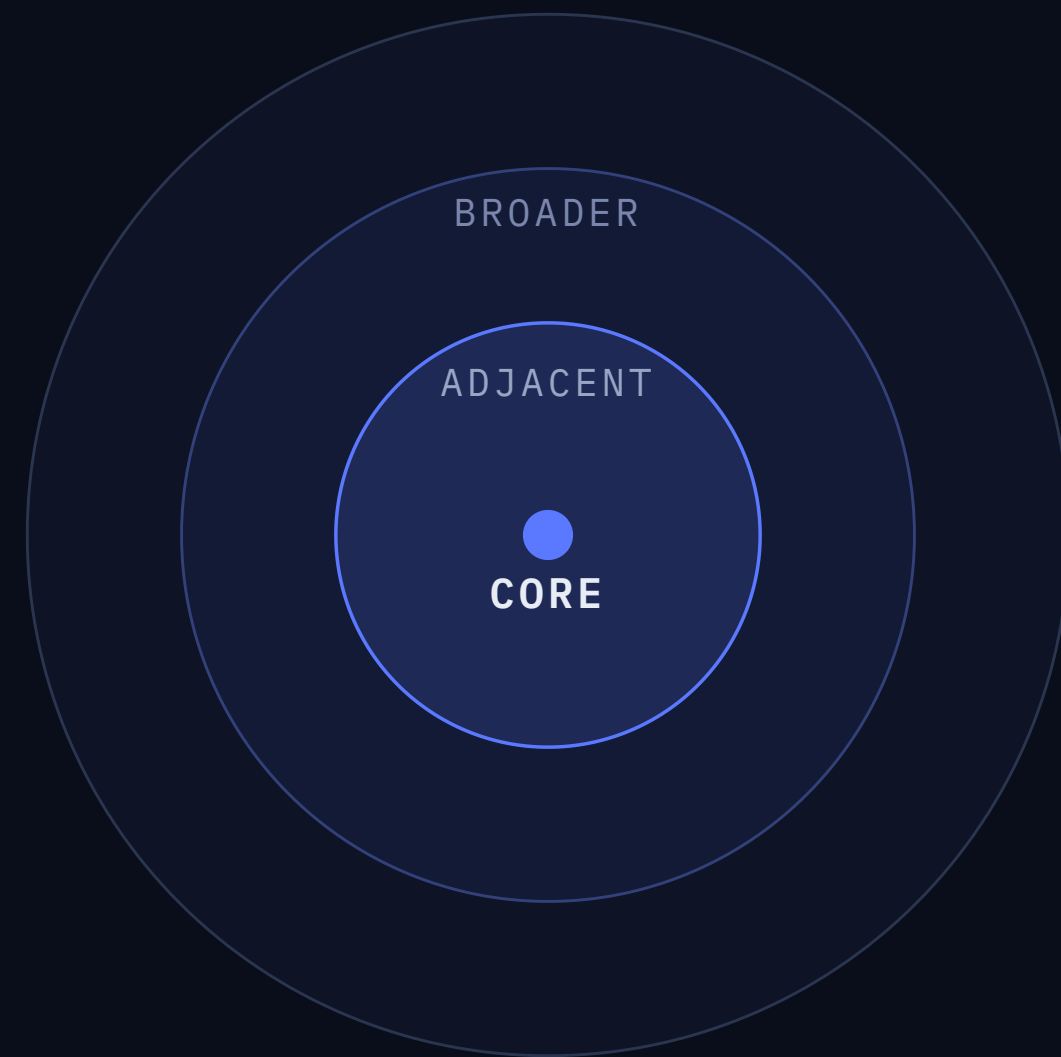
pay-per-screening; verdict reused at zero marginal cost

## FLYWHEEL

every screening makes the registry more valuable to the next agent



# From useful to mandatory.



## **CORE**

Every autonomous trading agent on Pharos — pre-flight safety is non-optional

## **ADJACENT**

Agent platforms, wallets, DeFi protocols — one call to integrate

## **BROADER**

The entire emerging AI-agent economy, across chains

As agents go from thousands to millions, a shared trust layer goes from useful to mandatory.

## THE MOAT

# Copy the code in an hour. Copy the filled registry — never.

### NETWORK EFFECT

Value = attestations × agents reading them. A fork starts empty. The data is the moat.

### STANDARD POSITION

Built Pharos-Skill-Engine native — the default trust call agents discover first.

### TRUST GRAPH

Every attestation is signed; reputation compounds to the earliest, most accurate attesters.

### NO-ADMIN CREDIBILITY

Ownerless + immutable = the neutral utility nobody must trust a company to use.

TRACTION

# Not a mockup — it runs.

|                                                     |                                                   |                                                       |                                                |
|-----------------------------------------------------|---------------------------------------------------|-------------------------------------------------------|------------------------------------------------|
| <div>Squid Game</div> <div>5</div> <div>BLOCK</div> | <div>SafeMoon</div> <div>48</div> <div>WARN</div> | <div>PancakeSwap</div> <div>88</div> <div>ALLOW</div> | <div>WBNB</div> <div>95</div> <div>ALLOW</div> |
|-----------------------------------------------------|---------------------------------------------------|-------------------------------------------------------|------------------------------------------------|

Live checker — paste any token, read the on-chain verdict. No backend, no database.

[agent-shield.vercel.app](https://agent-shield.vercel.app) ↗



ROADMAP

# From skill to autonomous guardian

• NOW

## Skill live

On-chain trust registry, deployed & verified on Pharos Atlantic.



• NEXT

## Sentinel

An autonomous agent that watches wallets, screens every token, refuses bad trades 24/7.



• LATER

## Trustless trust

Reputation-weighted + staked attestations.



# Agent Shield

The immune system for the autonomous agent economy.

X [x.com/agentshieldhq](https://x.com/agentshieldhq)

LIVE [agent-shield.vercel.app](https://agent-shield.vercel.app)

GITHUB [github.com/aaravfractal/agent-shield-pharos](https://github.com/aaravfractal/agent-shield-pharos)

CONTRACT [0x5ad1d58B95bb0FaF03499beF303788c73ba8aCe7](https://etherscan.io/address/0x5ad1d58B95bb0FaF03499beF303788c73ba8aCe7) · [verified](#) ·  
Pharos Atlantic (688689)

Built by [aaravfractal](#) · FAR AWAY 2026 · Agentic & Autonomous Systems